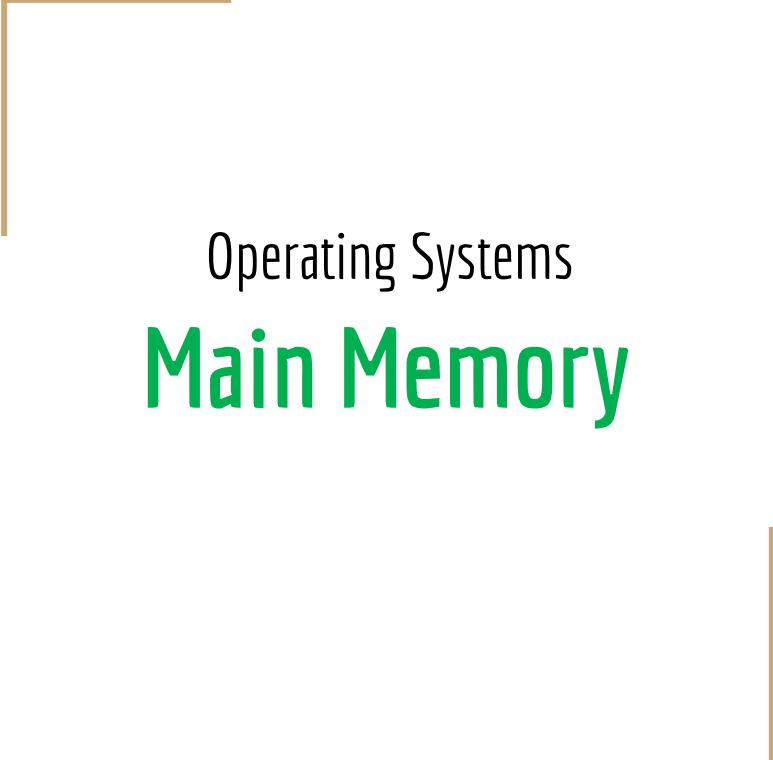


Operating Systems

Memory Management



Operating Systems

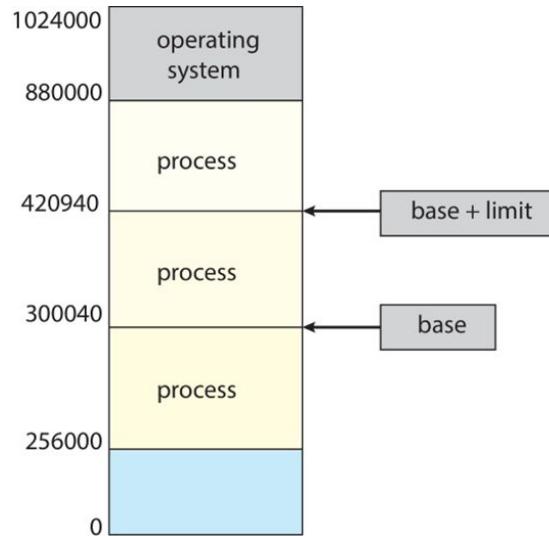
Main Memory

Background

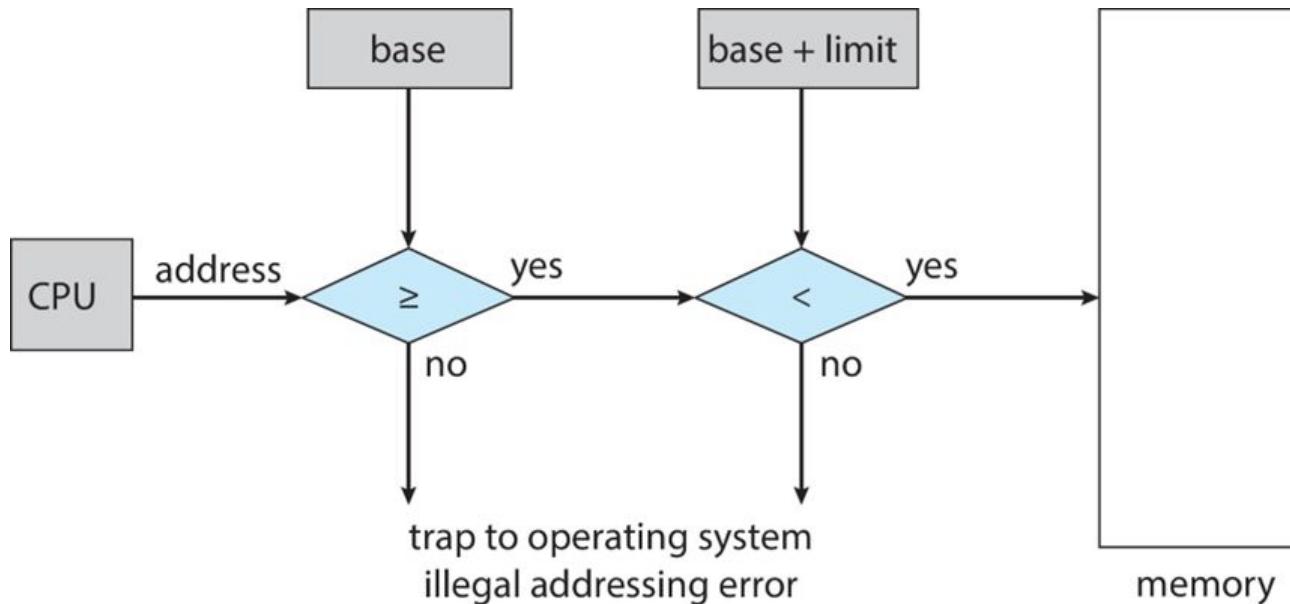
- ❑ Program must be brought (from disk) into memory and placed within a process for it to be run
- ❑ Main memory and registers are only storage CPU can access directly
- ❑ Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- ❑ Register access in one CPU clock (or less)
- ❑ Main memory can take many cycles, causing a **stall**, since it does not have the data required to complete the instruction that it is executing
- ❑ **Cache** sits between main memory and CPU registers for fast access
- ❑ Protection of memory required to ensure correct operation

Base and Limit Registers

- A pair of **base** and **limit registers** define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user



Hardware Address Protection

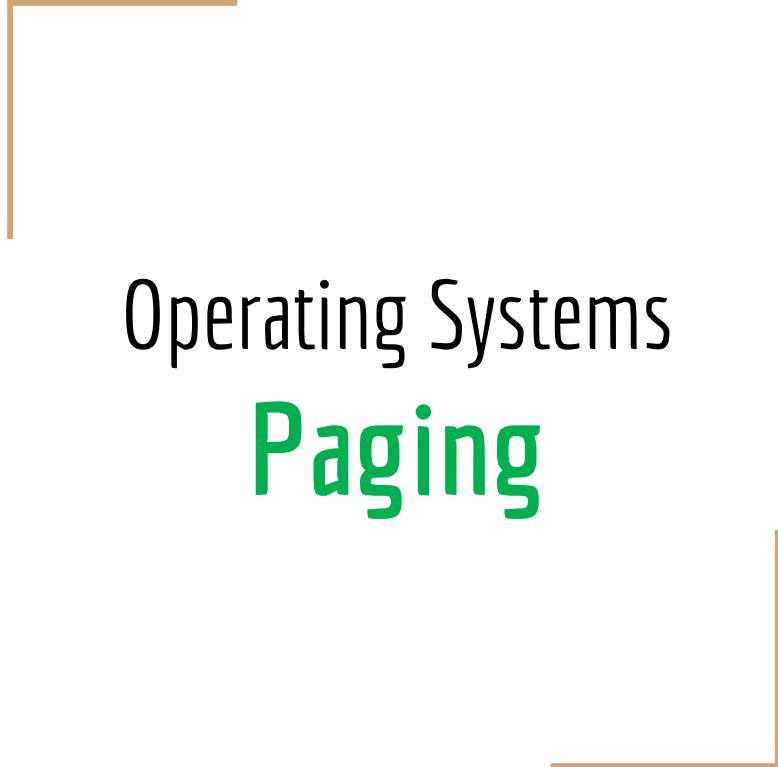




Operating Systems **Address Space**

Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program



Operating Systems

Paging

Paging

- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids external fragmentation
 - Avoids problem of varying sized memory chunks
- Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size N pages, need to find N free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Backing store likewise split into pages
- Still have Internal fragmentation

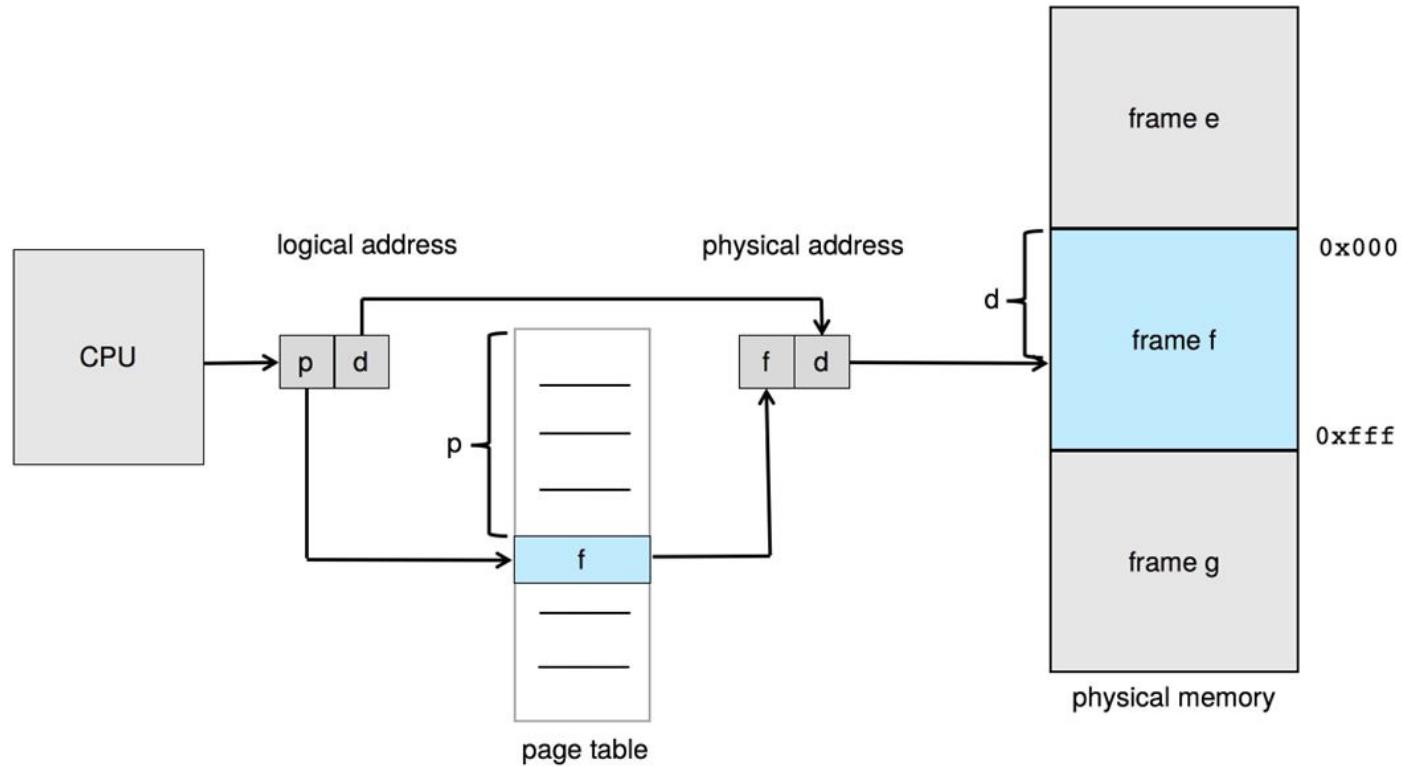
Address Translation Scheme

- Address generated by CPU is divided into:
 - **Page number (*p*)** – used as an index into a **page table** which contains base address of each page in physical memory
 - **Page offset (*d*)** – combined with base address to define the physical memory address that is sent to the memory unit

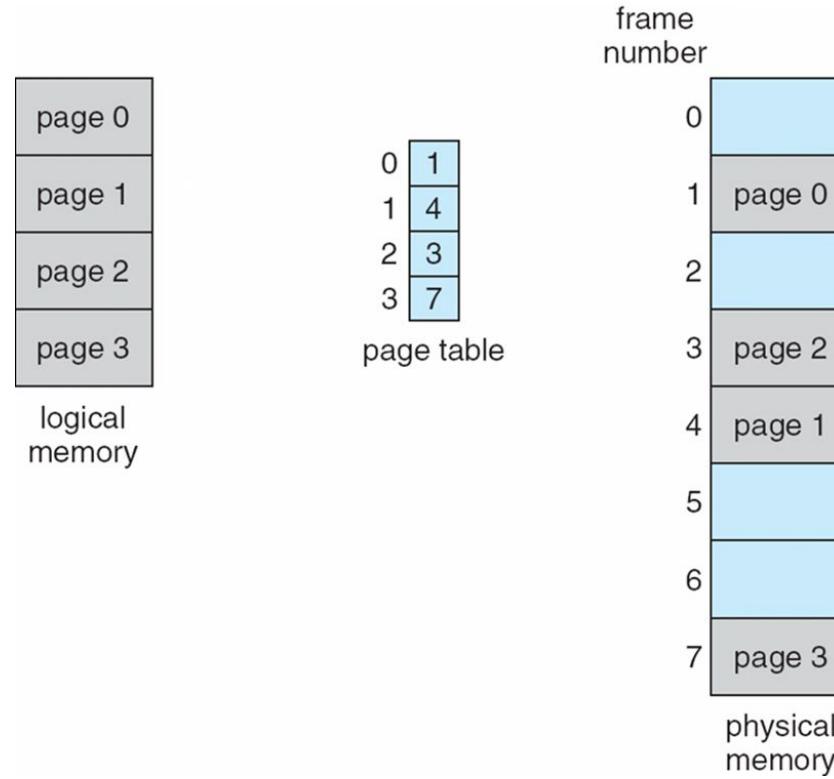
page number	page offset
<i>p</i>	<i>d</i>
$m - n$	n

- For given logical address space 2^m and page size 2^n

Paging Hardware



Paging Model of Logical and Physical Memory



Paging Example

$n=2$ and $m=4$

32-byte memory
and 4-byte pages

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

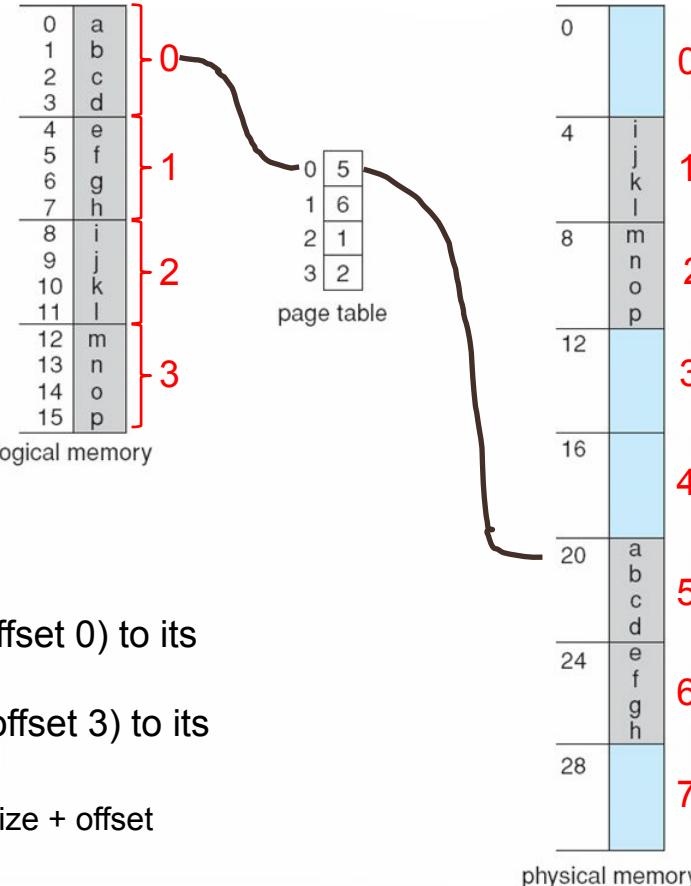
$$32/4 = 8 \text{ frames}$$

Paging Example

- m is used to determine the number of logical addresses, $2^m = 16$
- n indicates the offset within each logical address

$$n=2 \text{ and } m=4$$

 Physical memory → 32-byte memory and 4-byte pages



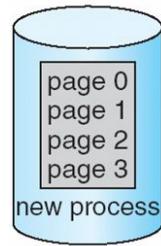
Size of Page

- Calculating internal fragmentation
 - Page size = 2,048 bytes
 - Process size = 72,766 bytes
 - 35 pages + 1,086 bytes
 - Internal fragmentation of $2,048 - 1,086 = 962$ bytes
 - Worst case fragmentation = 1 frame – 1 byte
 - On average fragmentation = $1 / 2$ frame size
 - So small frame sizes desirable?
 - But each page table entry takes memory to track
 - Page sizes growing over time
 - Solaris supports two page sizes – 8 KB and 4 MB
- Process view and physical memory now very different
- By implementation process can only access its own memory

Free Frames

free-frame list

14
13
18
20
15



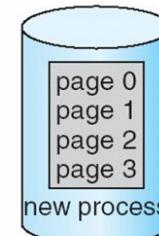
13
14
15
16
17
18
19
20
21

(a)

Before allocation

free-frame list

15



0 14
1 13
2 18
3 20

new-process page table

(b)

After allocation

free-frame list

13

14
page 0

15

16
page 1

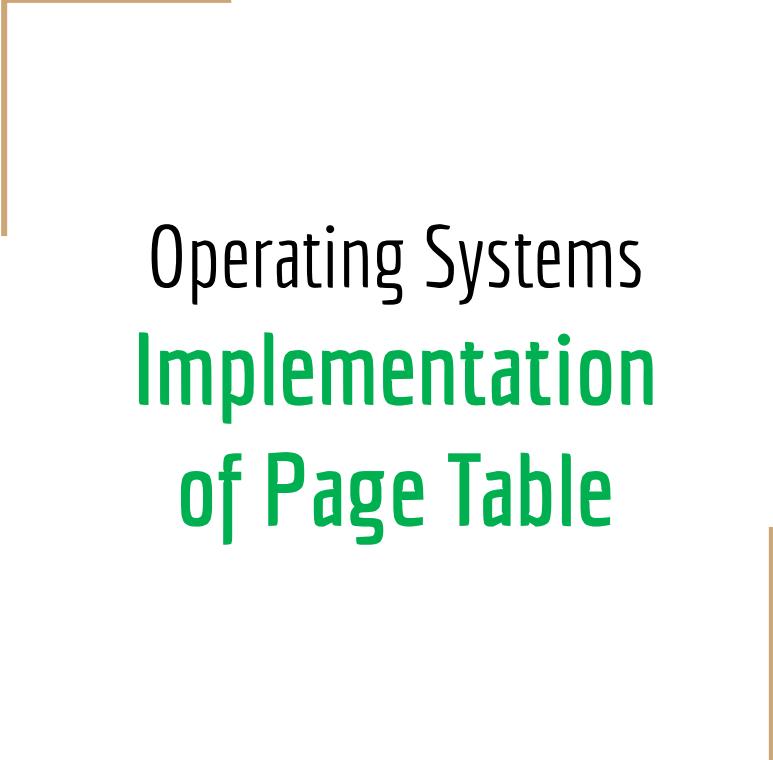
17

18
page 2

19

20
page 3

21



Operating Systems Implementation of Page Table

Implementation of Page Table

- Page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation lookaside buffers (TLBs)**

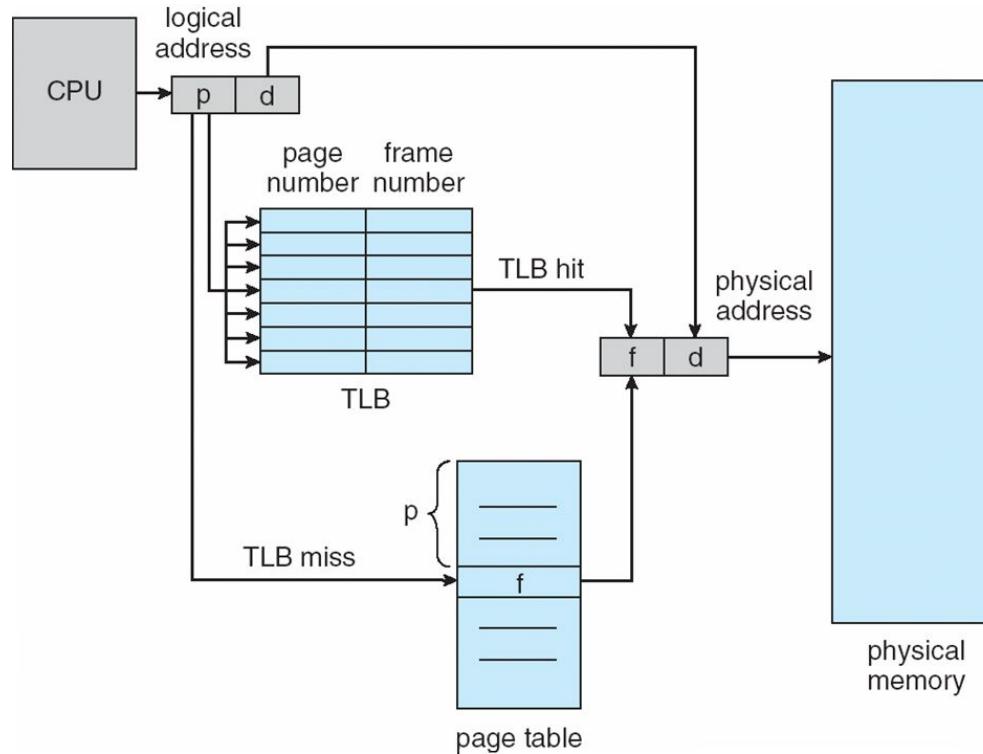
Associative Memory

- Associative memory – parallel search

Page #	Frame #

- Address translation (p, d)
 - If p is in associative register, get frame # out
 - Otherwise get frame # from page table in memory

Paging Hardware With TLB



Effective Access Time

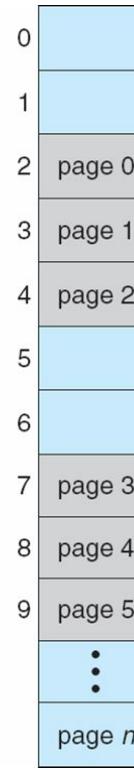
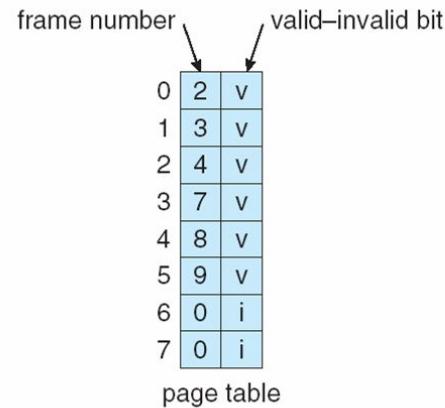
- Associative Lookup = ϵ time unit
 - Can be < 10% of memory access time
- Hit ratio = α
 - Hit ratio – percentage of times that a page number is found in the associative registers;
- **Effective Access Time (EAT):**
- Consider $\alpha = 80\%$, $\epsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $EAT = 0.80 \times 120 + 0.20 \times 220 = 140\text{ns}$
- Consider more realistic hit ratio -> $\alpha = 99\%$, $\epsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $EAT = 0.99 \times 120 + 0.01 \times 220 = 121\text{ns}$

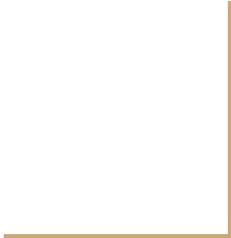
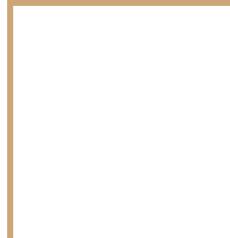
Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - “invalid” indicates that the page is not in the process’ logical address space
 - Or use **page-table length register (PTLR)**
- Any violations result in a trap to the kernel

Valid (v) or Invalid (i) Bit In A Page Table

00000	page 0
	page 1
	page 2
	page 3
	page 4
10,468	page 5
12,287	





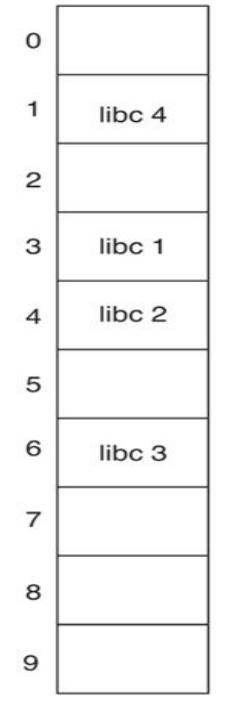
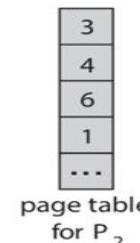
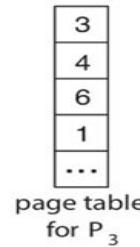
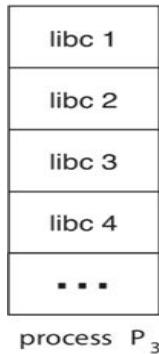
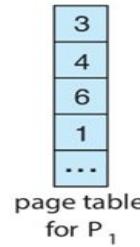
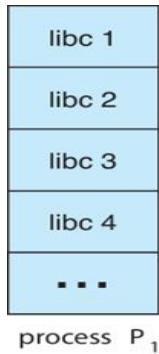
Operating Systems

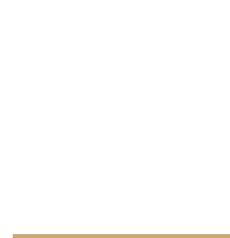
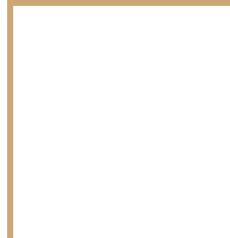
Shared Pages

Shared Pages

- **Shared code**
 - One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
 - Similar to multiple threads sharing the same process space
 - Also useful for inter-process communication if sharing of read-write pages is allowed
- **Private code and data**
 - Each process keeps a separate copy of the code and data
 - The pages for the private code and data can appear anywhere in the logical address space

Shared Pages Example





Operating Systems

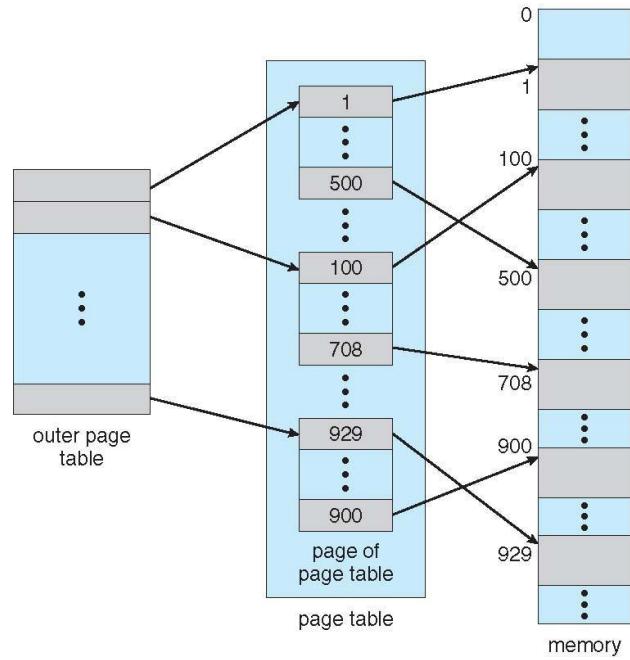
Hierarchical Page Table

Structure of the Page Table

- Memory structures for paging can get huge using straightforward methods
 - Consider a 32-bit logical address space as on modern computers
 - Page size of 4 KB (2^{12})
 - Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
 - That amount of memory used to cost a lot
 - Don't want to allocate that contiguously in main memory

Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table

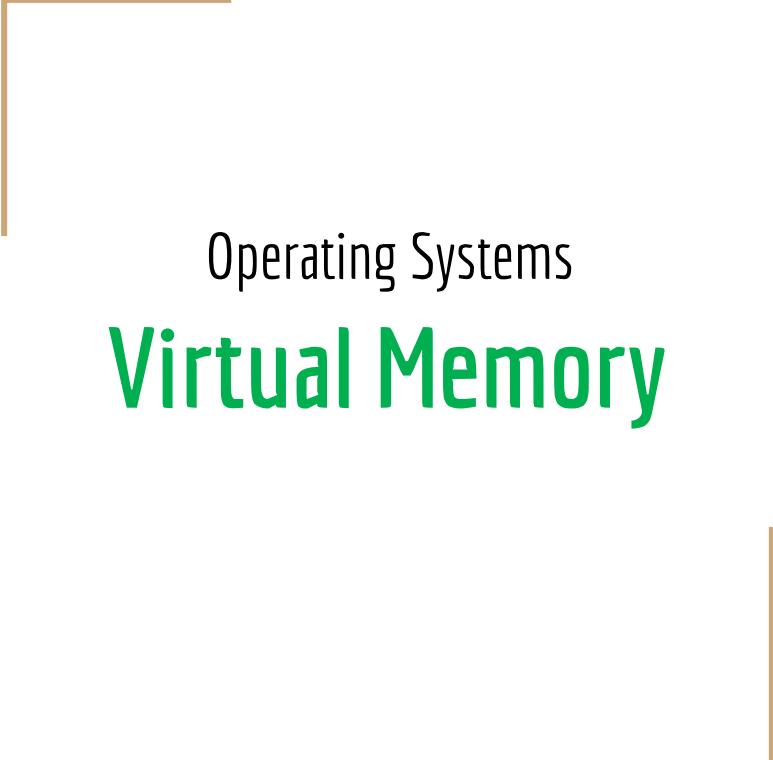


Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:
 - a page number consisting of 22 bits
 - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
 - a 12-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:

page number	page offset
p_1	p_2
12	10

- where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table
- Known as **forward-mapped page table**



Operating Systems

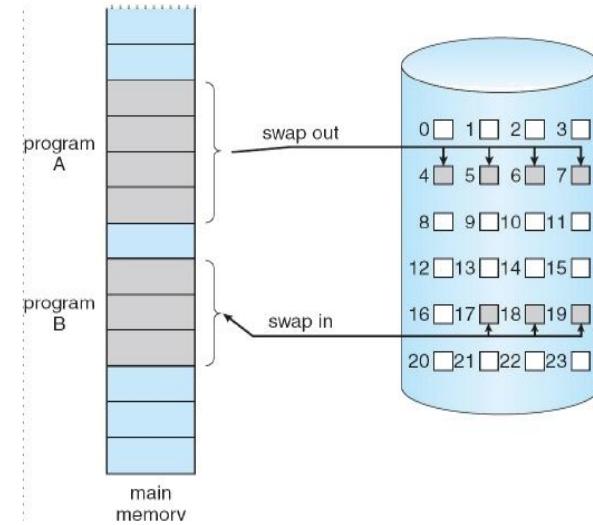
Virtual Memory

Background

- ❑ The term “virtual memory” refers to something which appears to be present but actually it is not.
- ❑ The virtual memory technique allows users to use more memory for a program than the real memory of a computer.
- ❑ Virtual memory is a **concept** that we use when we have processes that exceed the main memory.
- ❑ When computer runs out of physical memory, it writes its requirement to the hard disc in a swap file as “virtual memory”.

Demand Paging

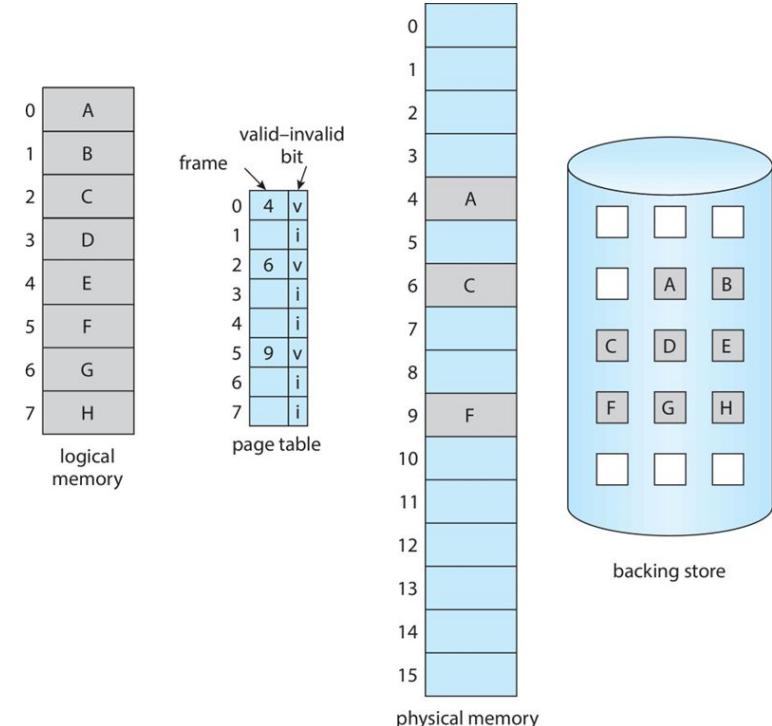
- ❑ Bring a page into memory only when it is needed
 - Less I/O needed
 - Less memory needed
 - Faster response
 - More users
- ❑ Disadvantage: Page fault interrupt
- ❑ Required hardware support:
 - Page Table with valid-invalid bit
 - Secondary memory



Demand Paging in OS

Valid-Invalid Bit

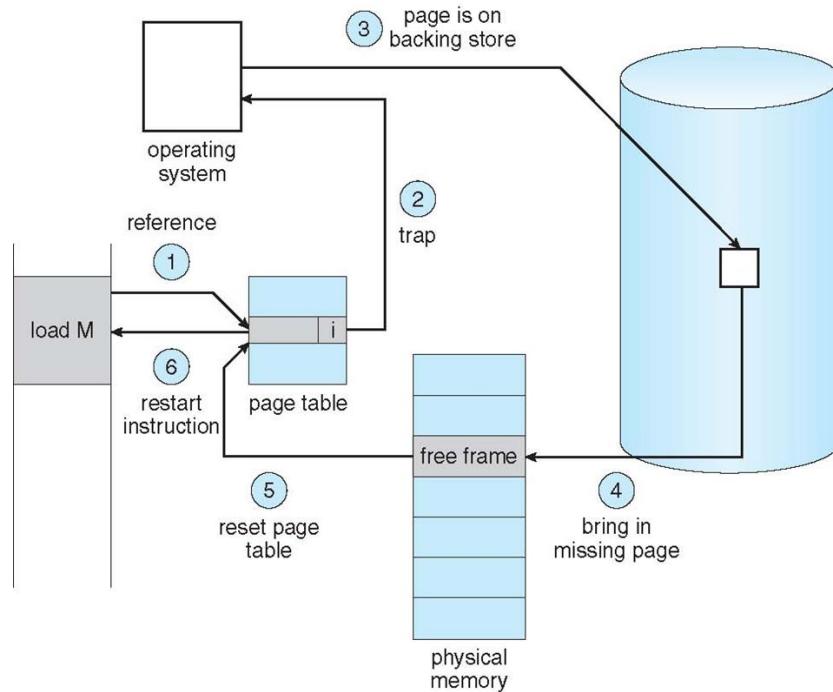
- ❑ An extra bit in the page table which indicates the existence of the page in the main memory.
- ❑ Attempt to access page
- ❑ If page is valid (in memory) then continue processing instruction as normal.
- ❑ If page is invalid then a page-fault trap / page-fault interrupt occurs.
- ❑ Page is needed \Rightarrow reference to it
 - Invalid reference \Rightarrow abort
 - Not-in-memory \Rightarrow bring to memory



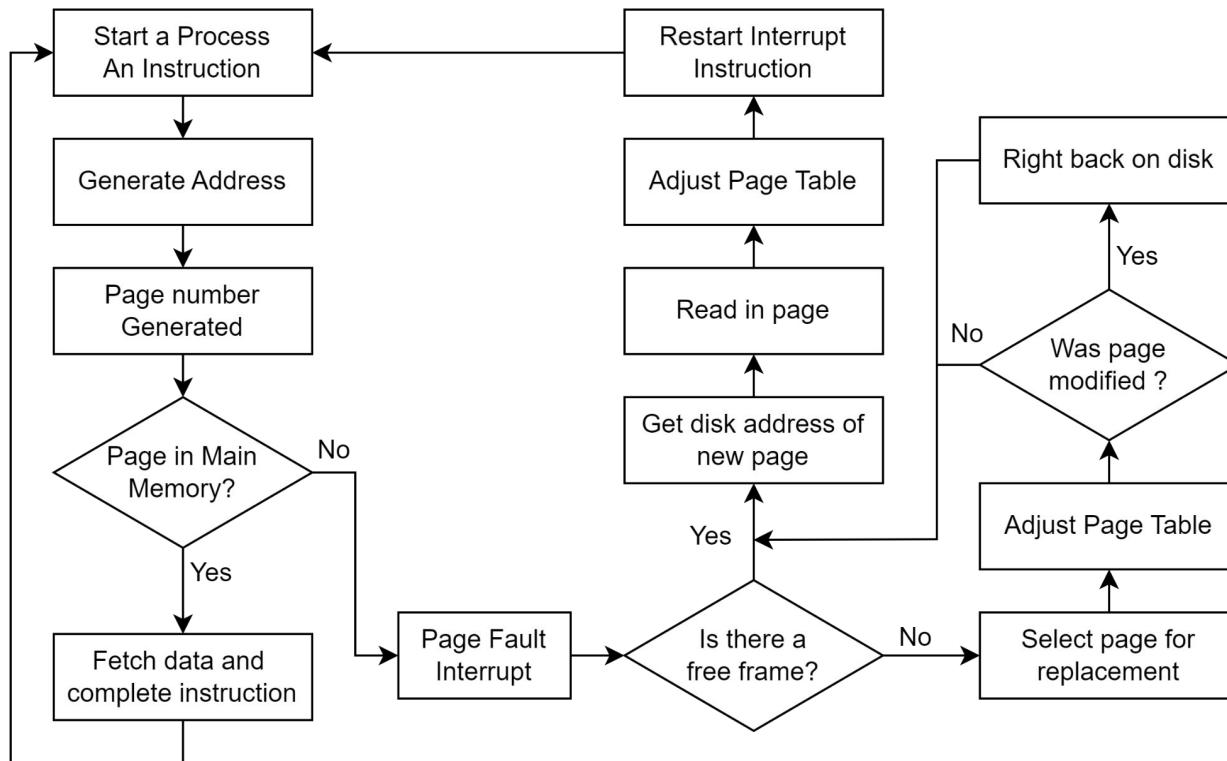
Page Fault

If there is ever a reference to a page, first reference will trap to OS \Rightarrow **page fault**

1. OS looks at another table to decide:
 - Invalid reference \Rightarrow abort.
 - Just not in memory.
2. Find empty/ free frame.
3. Load page from disk into frame.
4. Reset tables, validation bit = 1.
5. Restart instruction that caused page fault



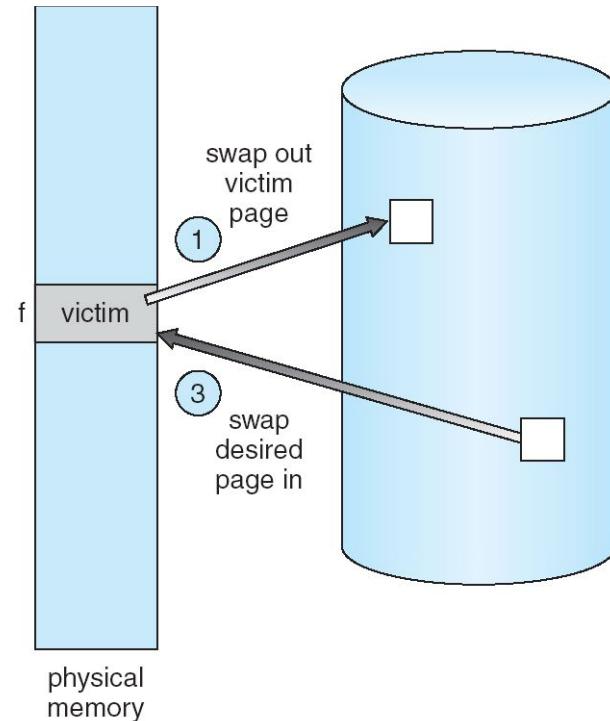
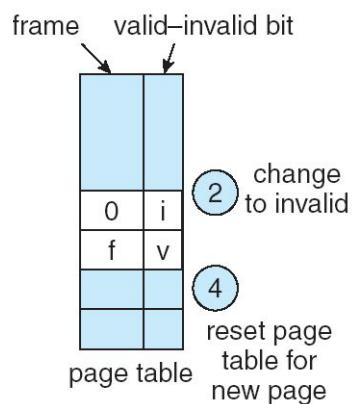
Demand Paging Flowchart



Page Replacement

Page Replacement Algorithms:

- FIFO (First In First Out)
 - LRU (Least Recently Used)
 - OPT (Optimal)



FIFO (First In First Out)

- Selects the page for replacement that has been in the memory for the longest amount of time

time	1	2	3	4	5	6	7	8	9	10	11	12
page	P2	p3	p2	p1	p5	p2	p4	p5	p3	p2	p5	p2
	p2*	p2*	p2*	p2*	P5	p5	p5*	p5*	P3	P3	P3	P3*
												
												
			hit					hit		hit		

Optimal

- ❑ Replace the page which is not used in longest dimension of time in future